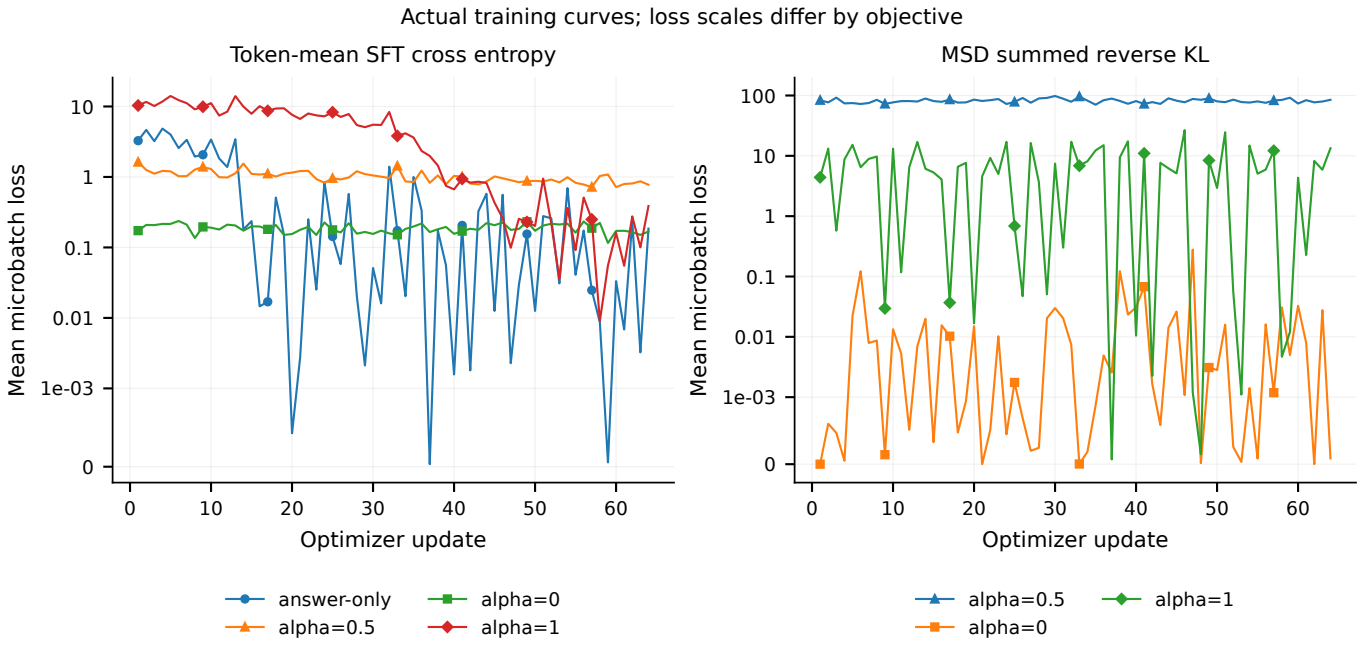
Actual training curves; loss scales differ by objective
Token-mean SFT cross entropy
MSD summed reverse KL
Mean microbatch loss
Optimizer update
10
1
0.1
0.01
1e-03
0
100
0
10
20
30
40
50
60
answer-only
alpha=0
alpha=0.5
alpha=1